

Predicting the Success Rate of Video Advertisements in AdTech

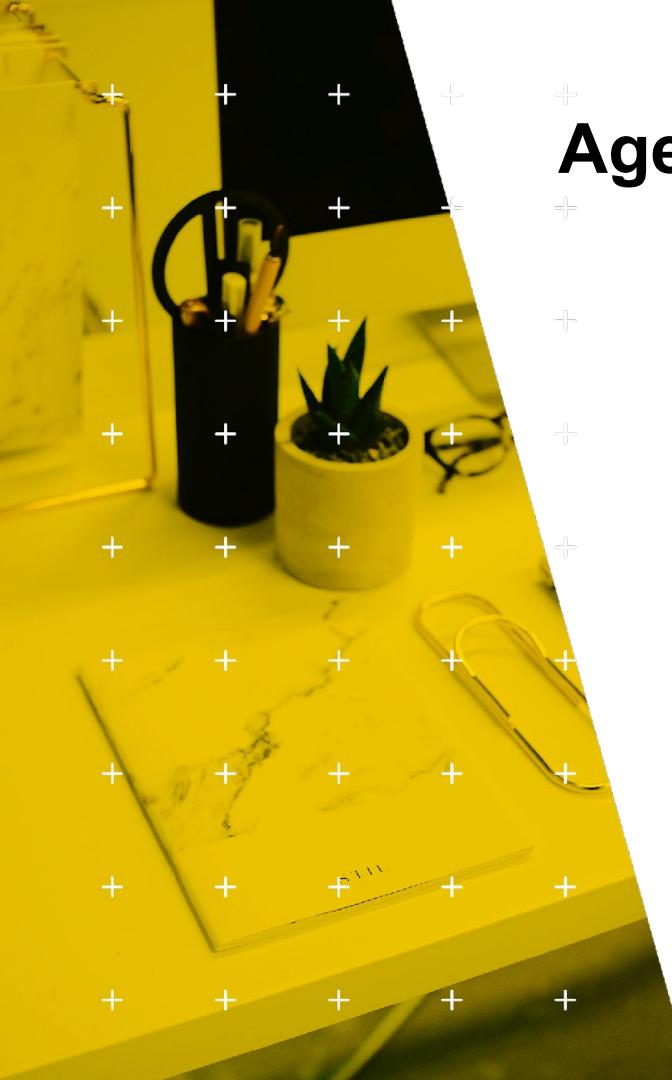
Buse Ozer

September 2020

Advisor: Boris Prodhomme

Tutor: Nacera Bennacer





Agenda

1. Introduction to Digital Advertising
2. General Overview of Advertising Ecosystem
3. Problem Definition
4. Methodology
5. Comparison of Different Models
6. Conclusion
7. Future Work



smart

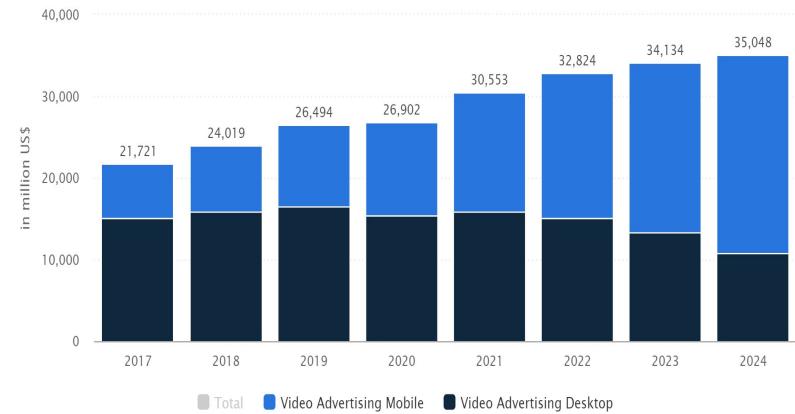
Introduction to Digital Advertising

Why Digital Display Advertising Is Important ?

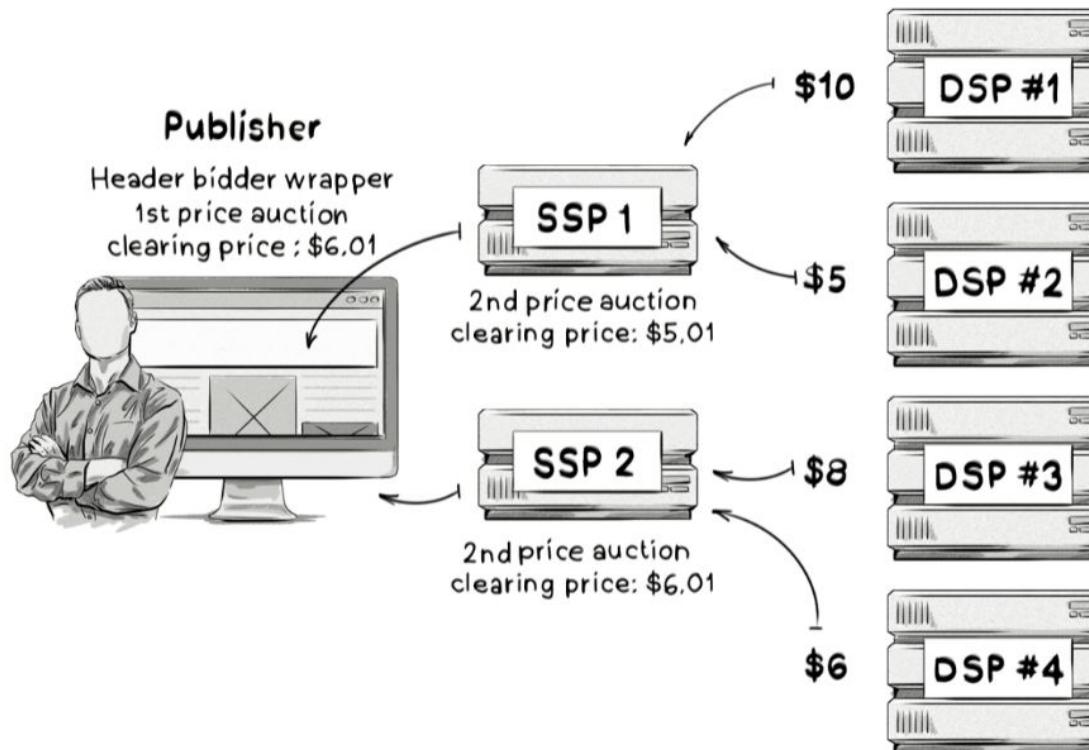
- High increase in digitalization
- High revenue in digital advertising
- Helps to accelerate digital commitment of users

Why Are Video Advertisements Important ?

- Drives more engagement
- Leads to more revenue
- Better user experience



General Overview of Advertising Ecosystem



Problem Definition

Success Rate Definition

Probability of an ad to be delivered

Success rate =
$$\frac{\text{\# of times a video ad displayed}}{\text{\# of times a video ad won an auction}}$$

Why Do We Focus on Video Success Rate ?



Our Display RTB delivery rate > 90 %



Our Video RTB delivery rate < 30%

Why Success Rate Optimization Is Important for Smart ?

Each loss opportunity is a direct loss in revenue for us and our publishers

smart

The Video Obstacles

Problems



- ❑ Timeouts of media files
- ❑ Player and creative incompatibility
- ❑ VPAID issues
- ❑ User leaves the page before the ad loads
- ❑ Connection cut

Solutions



- ❑ Ensure that video ads run smoothly
- ❑ Predict the delivery rate of a bid
- ❑ Highest expected revenue

Bidder	Bid price	Delivery rate (predicted)	Expected revenue
A	\$10	10%	?
B	\$4	50%	?

smart

The Objective

Increase of Smart and Publisher's revenue!



1

Apply delivery
rate optimization

2

Increase the delivery
rate of publishers

3

Increase the revenue
of publishers

4

Increase the revenue
of Smart

Predictive Model

What exactly is to be predicted ?

The probability of an ad to be displayed to an end-user

What is the goal of the predictor ?

1 Be as **accurate** as possible

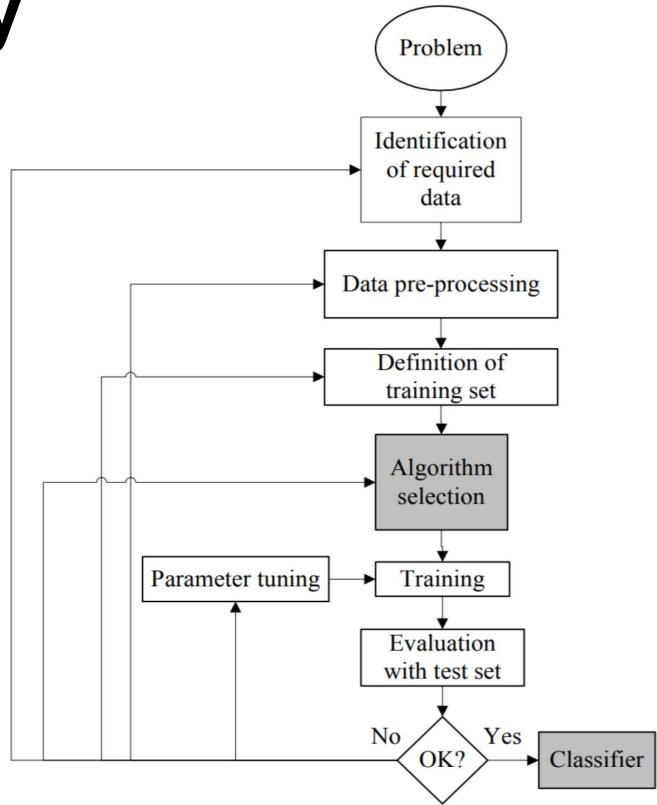
2 Have as **a few classification errors** as possible

What is the metric to measure the model performance ?

AUC, ability of a classifier to distinguish between classes

Among all “positive”-“negative” pairs in the dataset compute the proportion of those which are ranked correctly by the evaluated classification algorithm.

Methodology

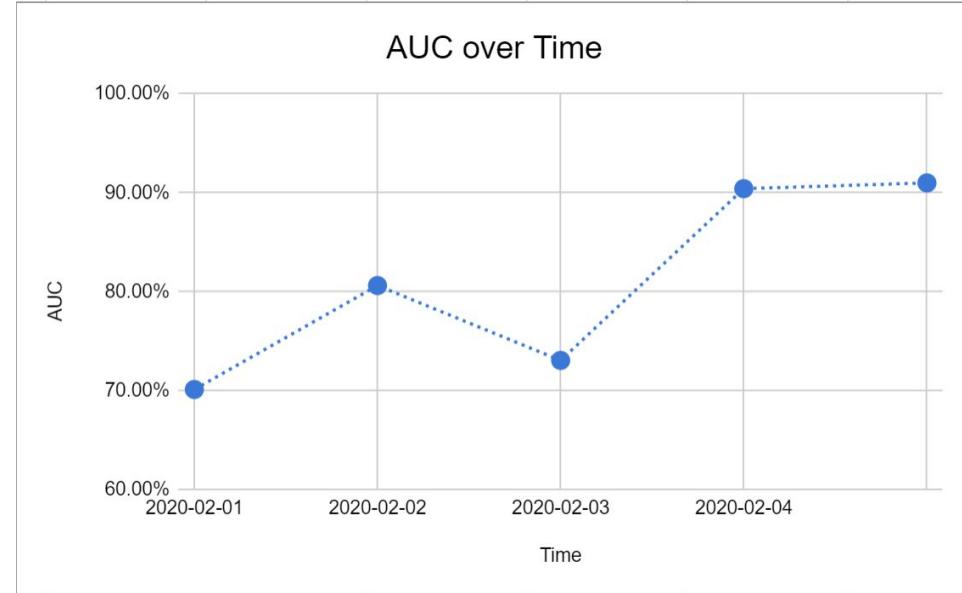


Data Set Split

DATA

Training Dev Test

Experiment	Training Set	Dev Set	Test Set
Experiment 1	Day 1	Day 2	Day 3
Experiment 2	Day 2	Day 3	Day 4
Experiment 3	Day 3	Day 4	Day 5
...			
Experiment n	Day n	Day n+1	Day n+2



Comparison of Tree-Based Models

Data Set	Model	LogLoss	Precision	Recall	Accuracy	AUC	F1
2020-02-01	Decision Tree	0.7477	0.8820	0.6223	0.6726	0.7942	0.7297
2020-02-01	Ada Boost	0.7259	0.8798	0.6099	0.6634	0.7705	0.7204
2020-02-01	Random Forest	0.6130	0.8789	0.6176	0.6678	0.7907	0.7254
2020-02-01	XGBClassifier	0.5960	0.8869	0.6242	0.6766	0.8059	0.7327
2020-02-02	Decision Tree	2.4462	0.7590	0.9019	0.7287	0.6389	0.8243
2020-02-02	Ada Boost	0.7950	0.7621	0.8882	0.7247	0.6562	0.8203
2020-02-02	Random Forest	0.4752	0.7496	0.9541	0.7452	0.7186	0.8396
2020-02-02	XGBClassifier	0.4716	0.7517	0.9358	0.7381	0.7330	0.8337
2020-02-03	Decision Tree	0.4975	0.8563	0.8996	0.8241	0.8728	0.8774
2020-02-03	Ada Boost	0.4750	0.8566	0.8960	0.8221	0.8431	0.8759
2020-02-03	Random Forest	0.3438	0.8548	0.9084	0.8285	0.8741	0.8808
2020-02-03	XGBClassifier	0.3081	0.8586	0.9062	0.8306	0.9042	0.8818



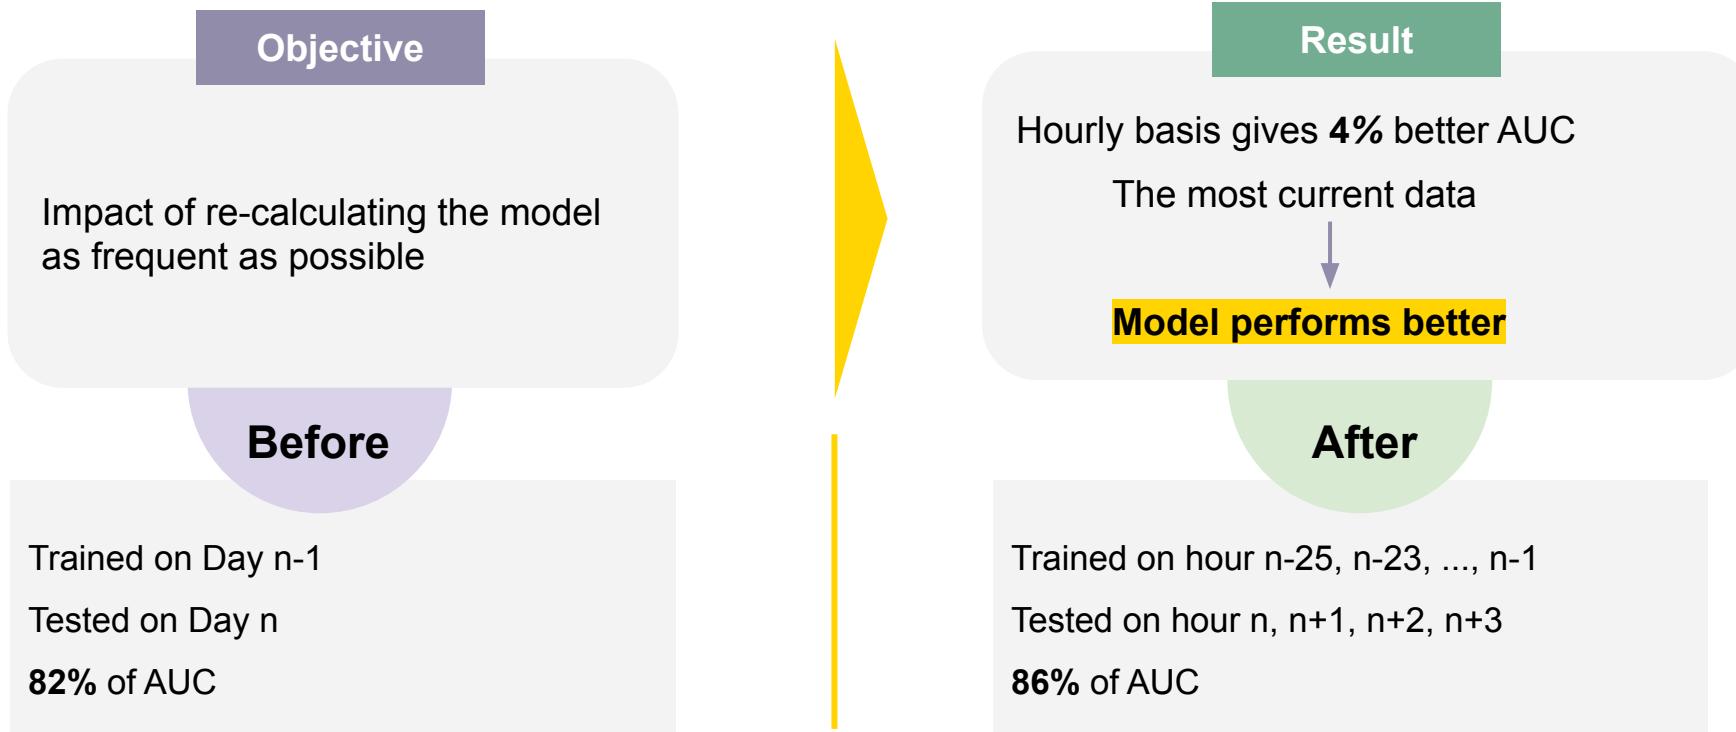
Comparison of Tree-Based Models

Model	Log Loss	AUC
Decision Tree	1.2305	0.7686
Ada Boost	0.6653	0.7566
Random Forest	0.4773	0.7945
XGBoost	0.4586	0.8144

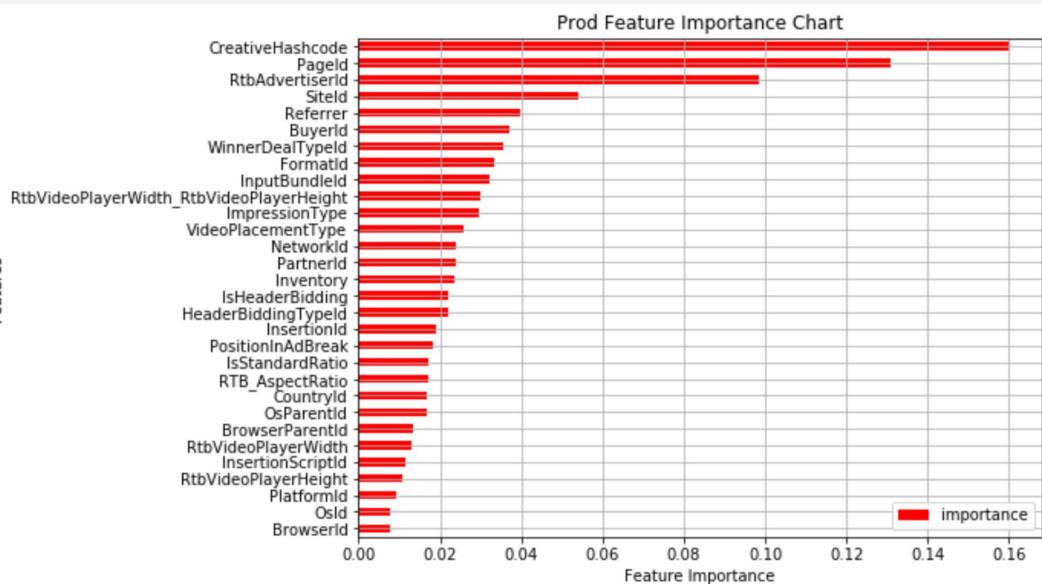
Averaged on Multiple Days

Model	Log Loss	Precision	Recall	Accuracy	AUC	F1
Decision Tree	1.2305	0.8324	0.8079	0.7418	0.7686	0.8105
Ada Boost	0.6653	0.8328	0.7980	0.7367	0.7566	0.8055
Random Forest	0.4773	0.8278	0.8267	0.7472	0.7945	0.8153
XGBClassifier	0.4586	0.8324	0.8221	0.7484	0.8144	0.8161

Importance of Freshness of The Model



Adding New Dimensions



New Features from RTB

3.21 %

PageId	OsParentId
BuyerId	OsId
WinnerDealTypeId	BrowserId
FormatId	RtbVideoPlayerWidth
CountryId	RtbVideoPlayerHeight
InsertionId	VideoPlacementType
BrowserParentId	PositionInAdBreak

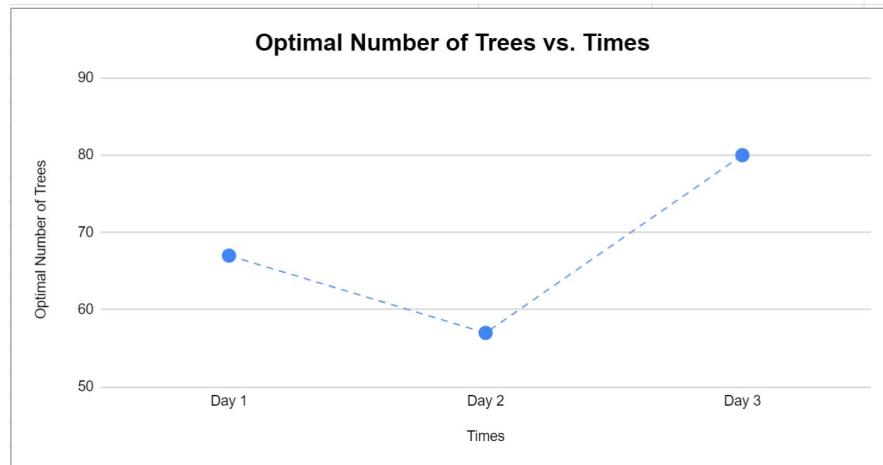
Early Stopping Strategy

Problem

There is a high variance between model complexity over time

Solution

Applying early stopping algorithm in order to control the model complexity



Early Stopping

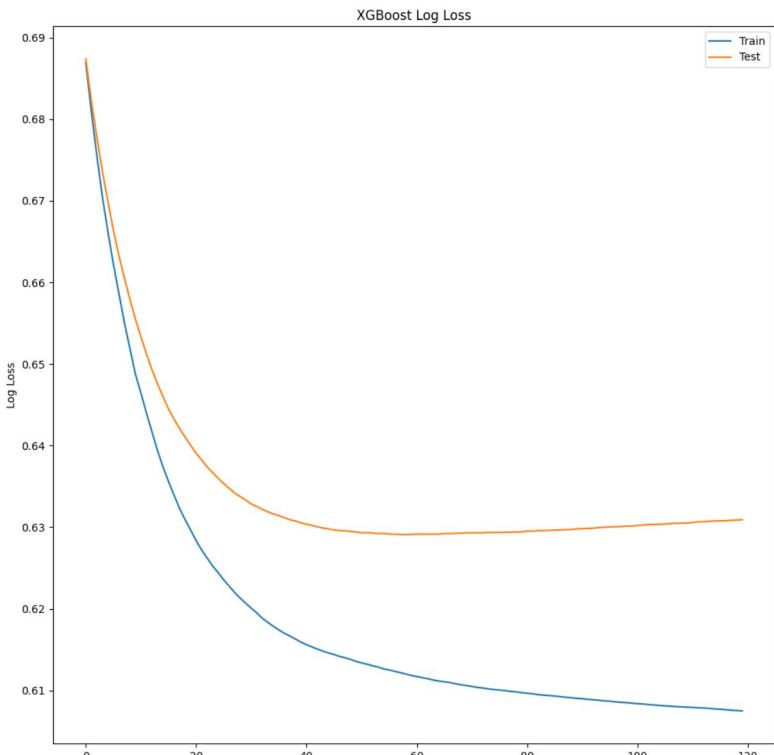
Use training and dev sets

Keeps on training the model until there is no more improvement on dev set

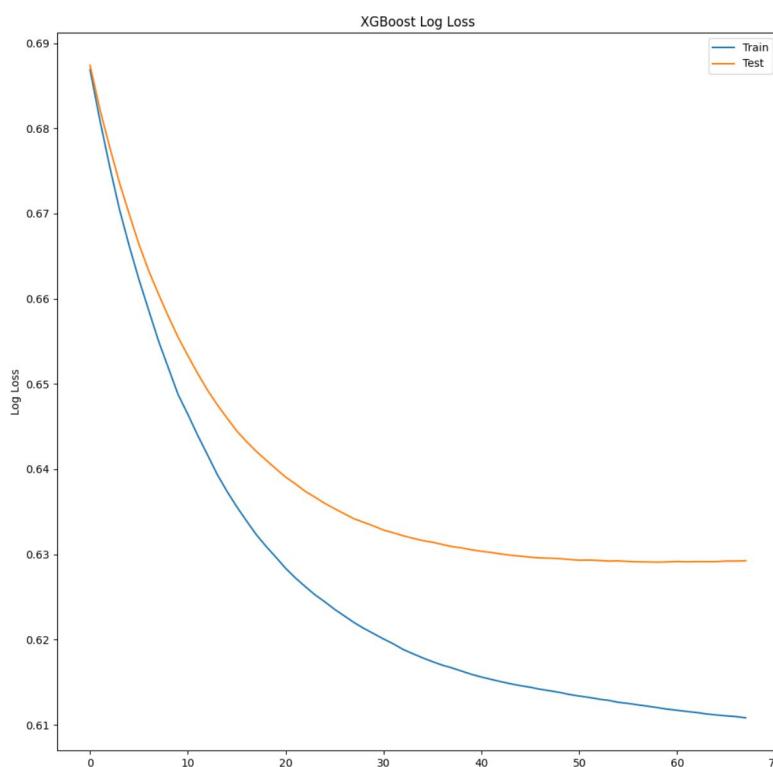
Determine how many trees to be built before overfitting

Prevents Overfitting

Without Early Stopping Algorithm



With Early Stopping Algorithm



Neural Network Model

Why Neural Network Model?

- Entity Embedding
- Ability to apply continuous learning
 - Ability of a model to learn continuously
 - Progressive and adaptive learning as new data comes in
 - Overcome forgetting of previously seen data
- Transfer Learning
 - The ability to transfer knowledge from one domain to another
 - Train a model on different tasks at the same time
- TensorFlow Serving in production

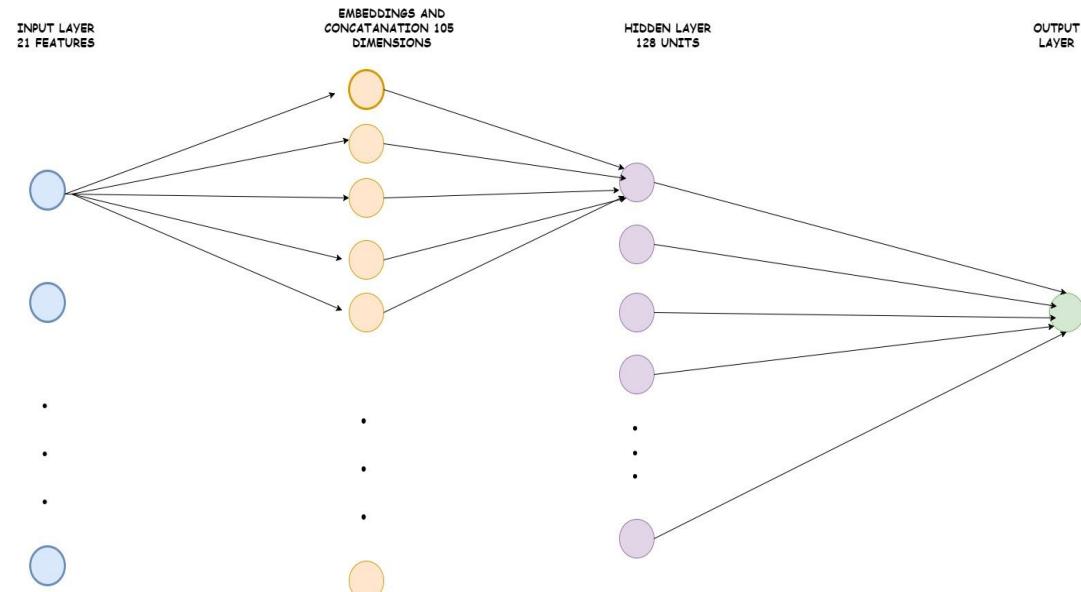


Neural Network Model

- Embedding layer with 5 dim
- 1 hidden layer with 128 units
- Dropout with 0.3 probability

 TensorFlow

 Keras



Final Results

XGBoost Model

Training
AUC
92.4%

Test AUC
89.8%

Neural Network Model

Training
AUC
91.5%

Test AUC
89.2%



Conclusion

- ❖ There is **no trend** in AUC over time
 - Run experiments on multiple days
- ❖ XGBoost model **outperforms** Decision Tree, AdaBoost and Random Forest
 - Better AUC, better **control over the model**
 - Scalable and parallelizable
- ❖ High variance in data
 - The model has to be **adaptive** and **flexible**
 - Early Stopping mechanism
- ❖ XGBoost and Neural Network performs nearly the same
 - Neural Network requires more expertise
- ❖ Implementation of Neural Network has some advantages
 - Continuous Learning
 - Serving in production
 - Transfer Learning

Future Work

- 
- ❖ Change the weighting strategy of the XGBoost model
 - Weighting on the revenue of samples to focus where the impact is high
 - ❖ Continuous Learning in Neural Network Model
 - ❖ Anomaly Detection
 - Detect spaces where delivery late is low
 - Finding the responsible publisher for the anomaly

Thank you !

Predicting the Success Rate of Video Advertisements

Buse Ozer

September 2020



References

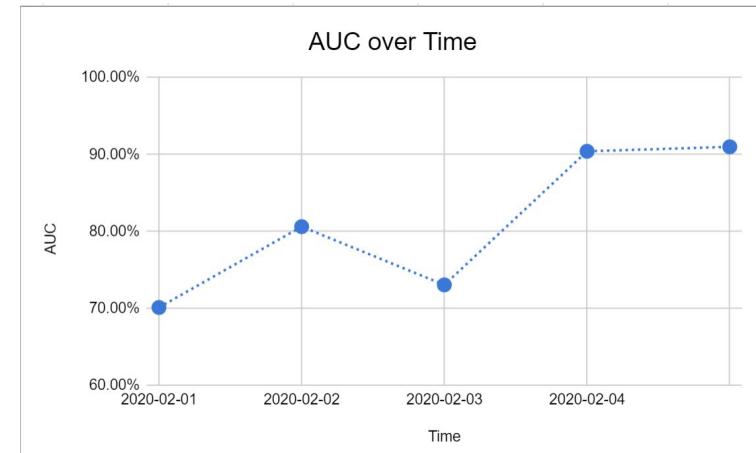
- ❖ Page 3,
<https://www.statista.com/outlook/218/100/video-advertising/worldwide>
- ❖ Page 4,
<https://clearcode.cc/blog/sequential-auctions-header-bidding-first-price-second-price-auctions/>
- ❖ Page 9, Osisanwo, F. Y., Akinsola, J. E. T., Awodele, O., Hinmikaiye, J. O., Olakanmi, O., & Akinjobi, J. (2017). Supervised machine learning algorithms: classification and comparison. International Journal of Computer Trends and Technology (IJCTT), 48(3), 128-138.

Q&A

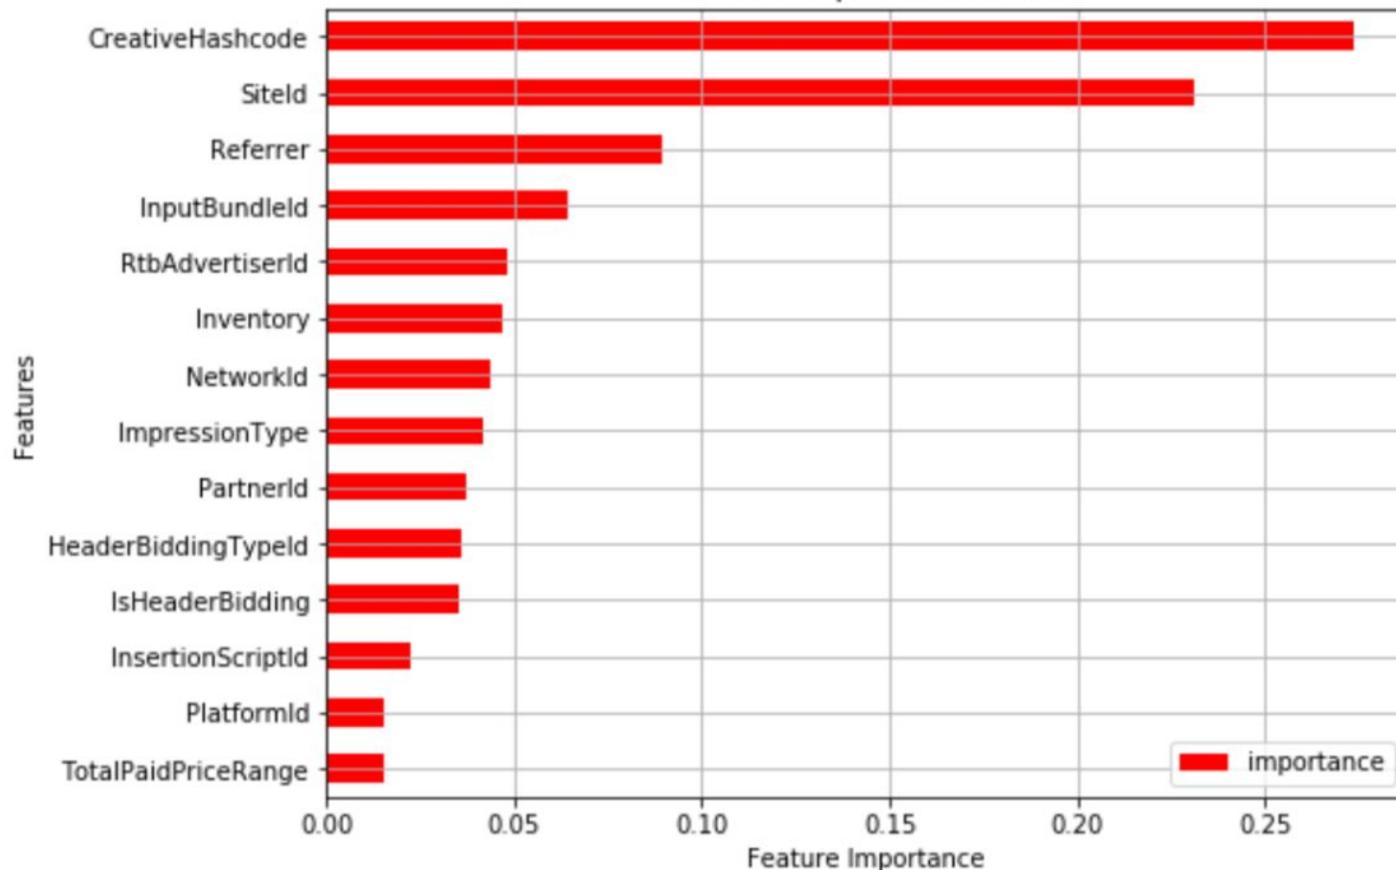


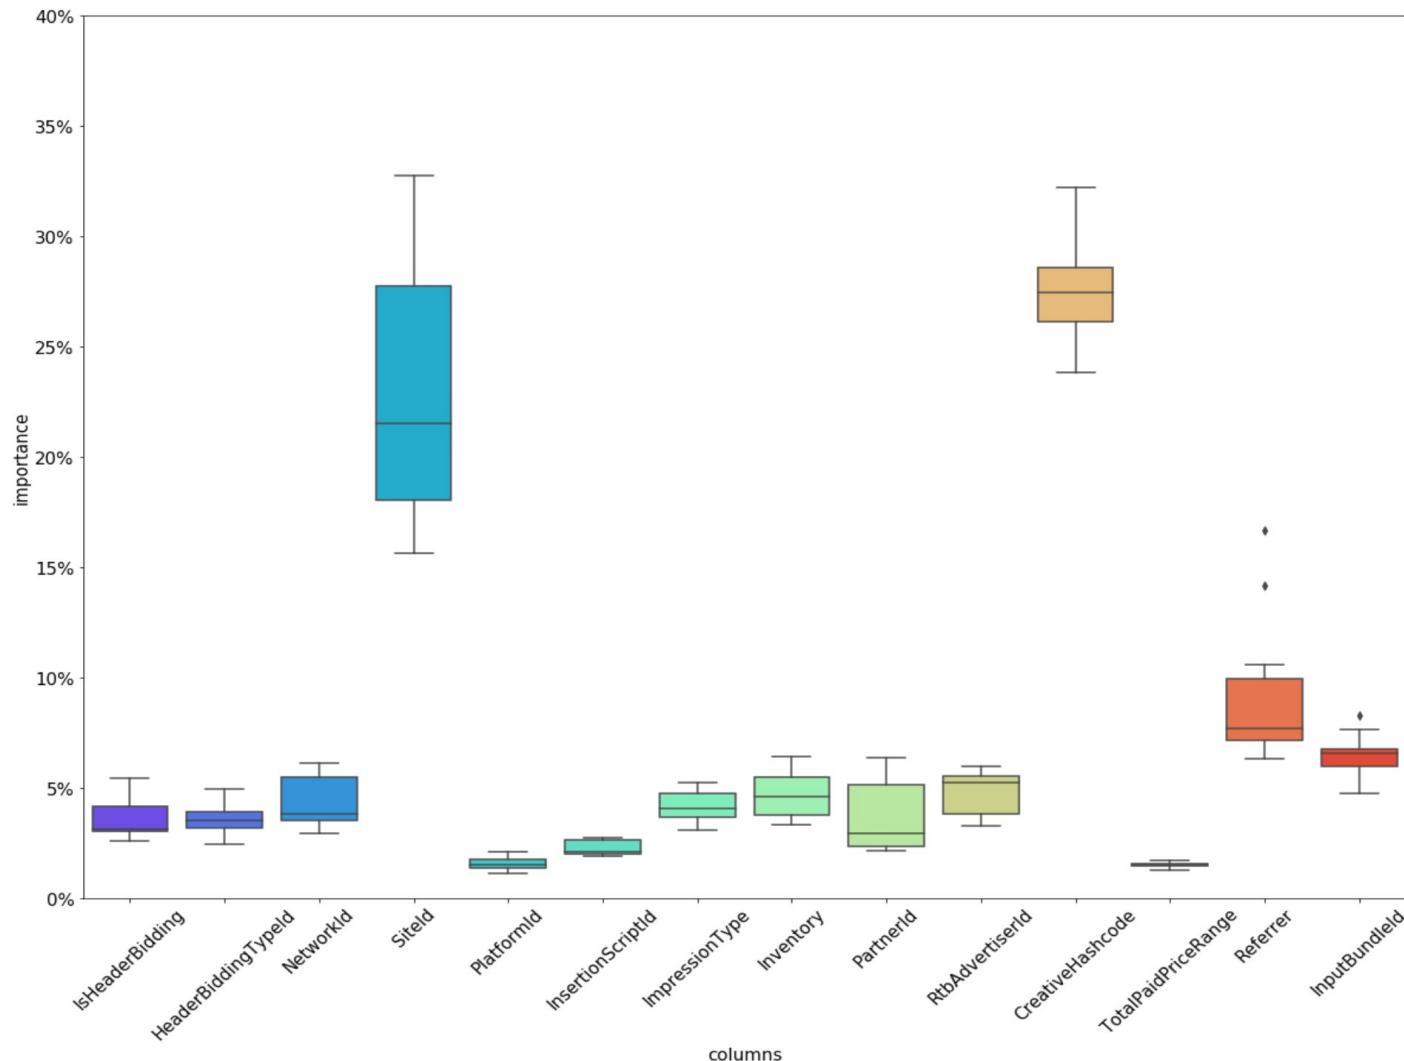
Tuning Hyperparameters of the Model

- ❑ Grid Search
- ❑ Tested on multiple days
- ❑ Best results are obtained with the parameters below
 - ❑ **Max_depth = 20**
 - ❑ Maximum tree depth for base learners
 - ❑ **Learning_rate = 0.05**
 - ❑ Weighting factor for the corrections from existing trees
 - ❑ **Colsample_bytree = 0.6**
 - ❑ The percentage of dimensions to built the trees
 - ❑ **N_estimators = 140**
 - ❑ Number of trees to use in XGBoost



Feature Importance Chart



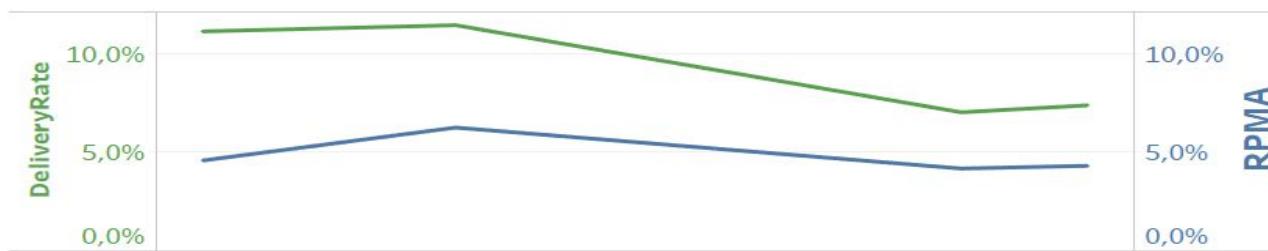


smart+

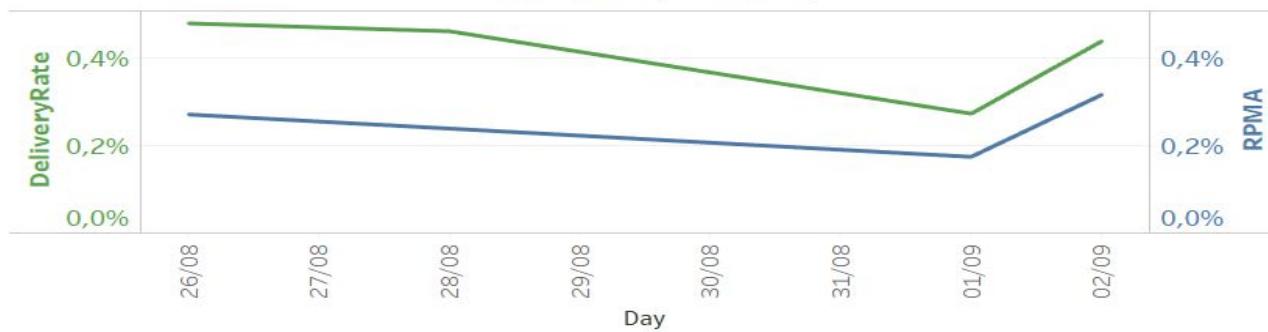
Daily revenue (Total & Optimizable)



Daily uplifts (on Optimizable)



Daily uplifts (on Global)



smart+

Only Video inventory

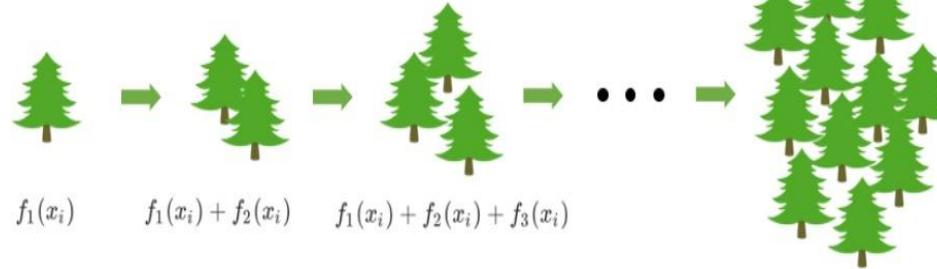
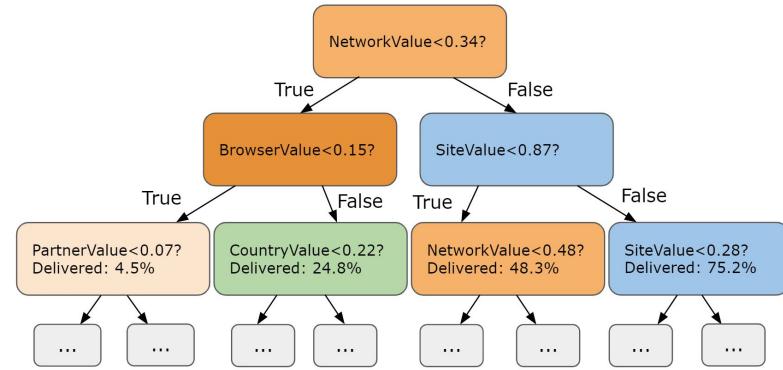
- Optimizable auctions = multi-bid auctions
- Optimized auctions = Optimizable and DROS was called
- Optimizable but not optimized = DROS was not called (for A/B testing)

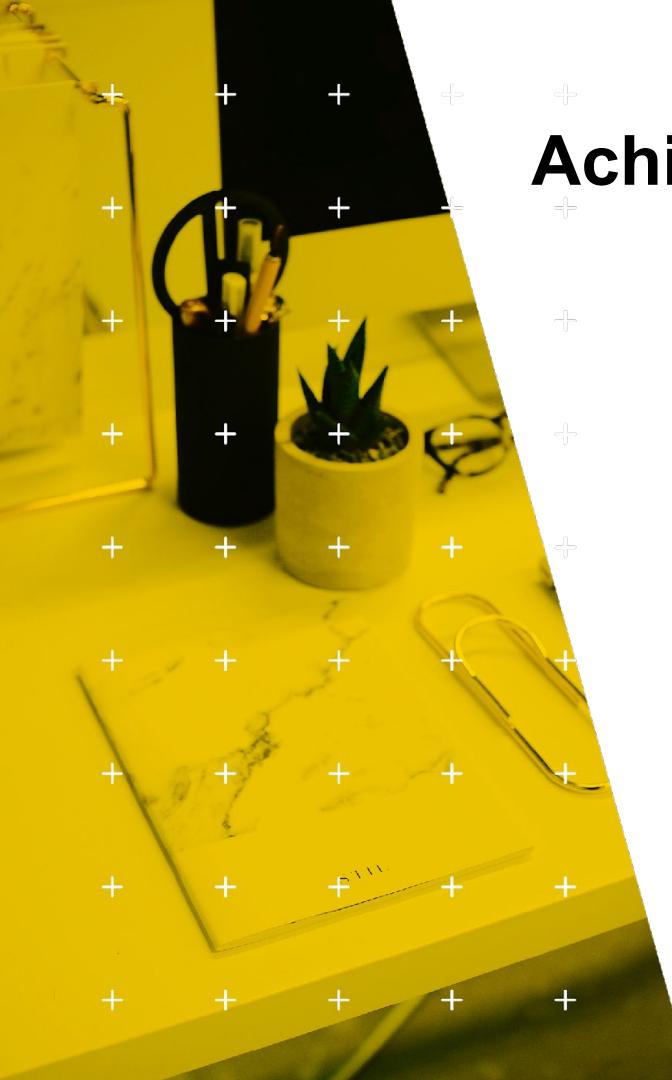
OptimizableRatio (Impressions)	4,8%
OptimizableRatio (Revenue)	6,2%
DeliveryRate (Global)	21,3%
DeliveryRate (Optimizable but not Optimized)	20,7%
DeliveryRate (Optimized)	22,5%
DeliveryRateUplift (on Optimizable) (pts)	1,9pts
DeliveryRateUplift (Global) (pts)	0,1pts
DeliveryRateUplift (on Optimizable) (%)	9,2%
DeliveryRateUplift (Global) (%)	0,4%
RevenuePerThousandWin (Global)	1,07 €
RevenuePerThousandWin (Optimizable but not Optimized)	1,33 €
RevenuePerThousandWin (Optimized)	1,39 €
DailyRevenueUplift (Global)	226 €
RevenueUplift (on Optimizable)	4,9%
RevenueUplift (Global)	0,3%

XGBoost Model

Model

- Tree-based
- Ensemble Learning
- Consists of many simple trees
- Learns mistakes from previous trees
- Builts sequential trees





Achievements

1. Moving all experiments to Google Cloud Platform
2. Importance of Freshness of the Model
3. Adding New RTB dimensions to the model
4. Analyzing Cross Features
5. Tuning Hyperparameters of the Model
6. Early Stopping Algorithm
7. Pushing the new model to production
8. PoC of a new model with Deep Learning

1

Moving all Experiments to Google Cloud Platform



Problems

- Resource limitation
- Time limitation
- Everything is in local machines

Solution



Google Cloud Platform



AI Platform



Big Query



GPU

How ?

- Allocate resources as much as needed
- Improved Performance
- Allows collaboration, accessible amongst team members
- Cloud Billing Account
- State of the art security

smart

Better Process for Evaluating the Model

In production

- Learn from past
- Predict on the present

Goal

- Reproduce “in-production” conditions

How ?

- Split historical dataset into
 - Training/Testing
- Time constraint

Data Set Split

DATA

Training

Dev

Test

Used during the training and tuning of the model

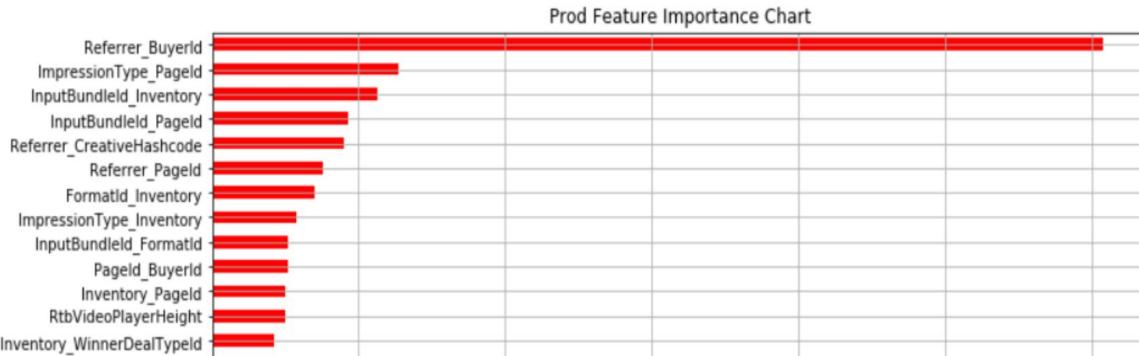
Don't touch this till the very end



4 Analyzing Cross Features

Objectives

- See the most important feature combinations
- Find incompatibilities between supply and demand dimensions
- Gain knowledge for anomaly detection

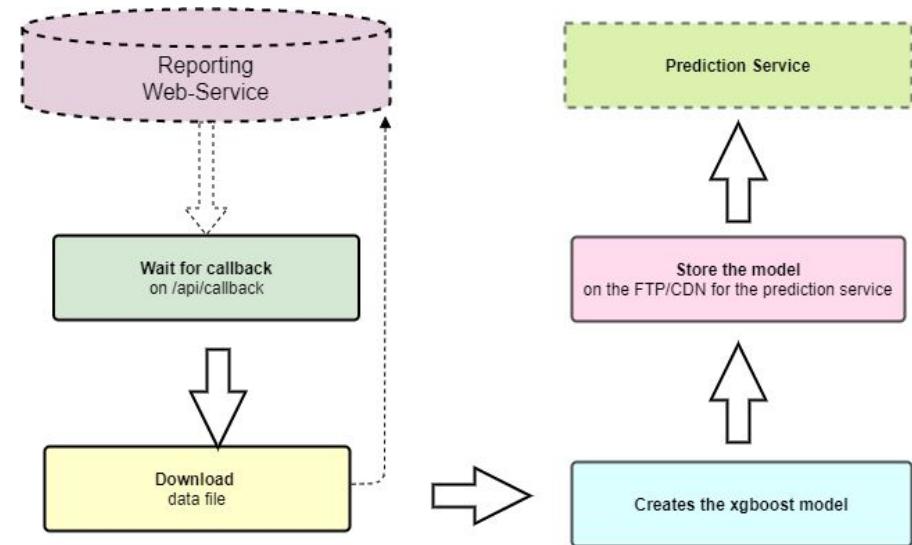


smart+

Pushing the new model to production

Delivery Rate Learning Service

- Deployed the model in production
- Scheduled a new reporting for model B
- Created new configuration for prod 2



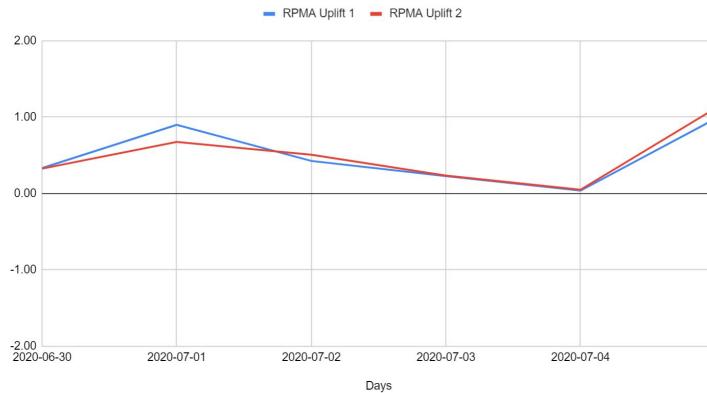
XGBoost



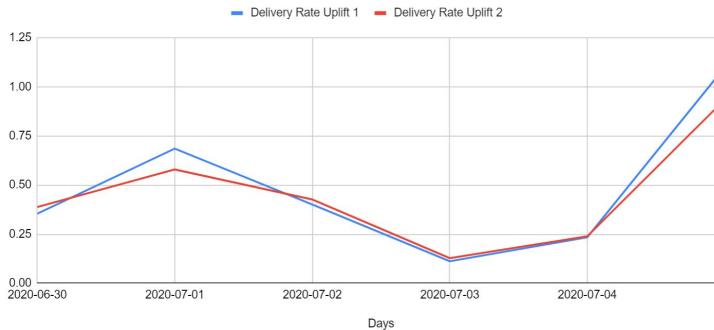
smart+

The First Results

RPMA Uplift for Model 1 and RPMA Uplift for Model 2



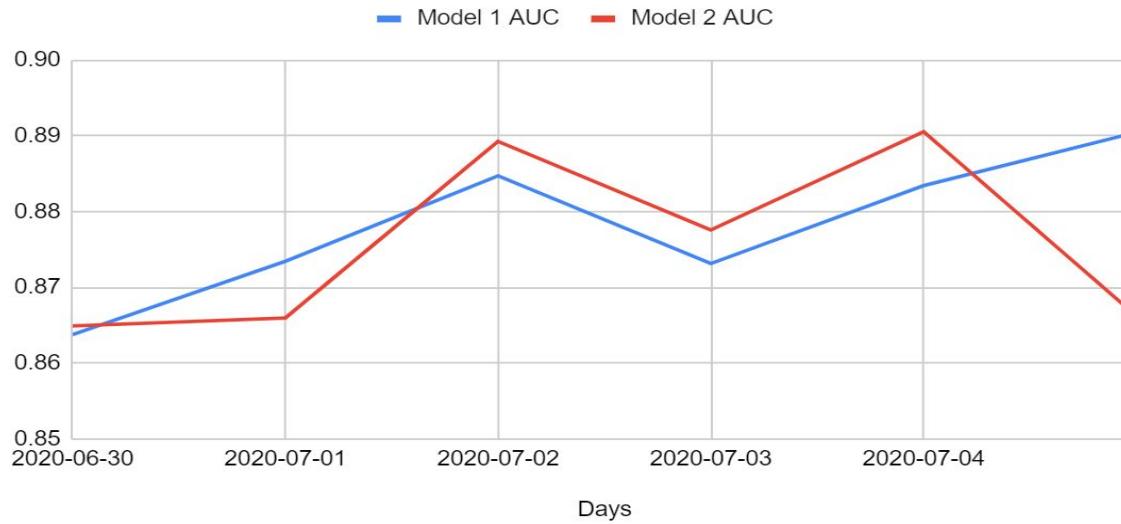
Delivery Rate Uplift for Model 1 vs Delivery Rate Uplift for Model 2



- 06-30
 - Activation of the new model, same configuration
- 07-01
 - Making experiments
 - Adding new dimensions
 - Changes of parameters of the model
 - Testing different conf on reporting scheduling
 - Model B is unstable
 - Learn less, less precise
- 07-02, 07-03, 07-04
 - No changes, model B is stable
 - Better precision
 - No uplift when there is no anomalies
- 07-05
 - Tests for memory leak
 - Model B is unstable
 - Learn less, less precise

The First Results

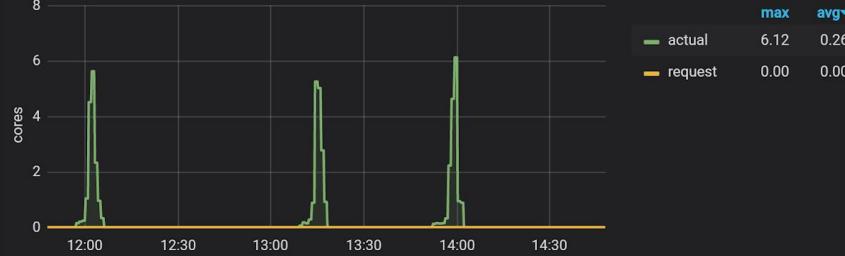
Model 1 AUC and vs Model 2 AUC



Model 1

Usage total

Cpu usage (total)



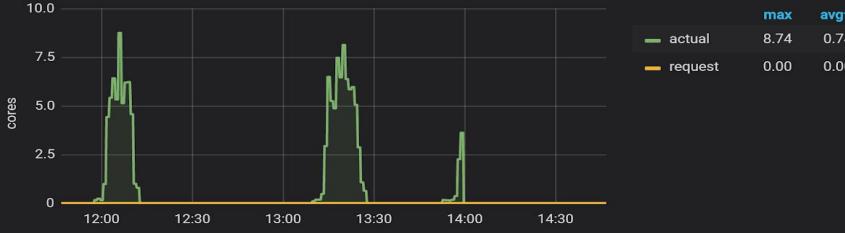
Memory usage (total)



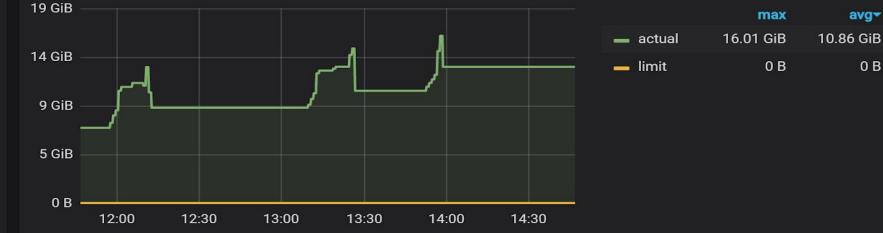
Model 2

Usage total

Cpu usage (total)



Memory usage (total)

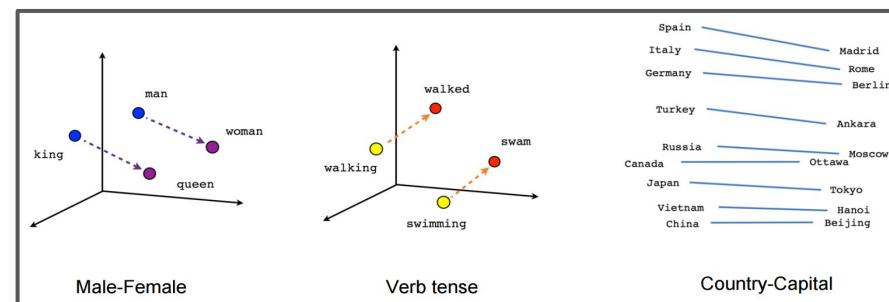


PoC of a new model with Deep Learning

Handle Categorical Data with Deep Learning

- ❑ One Hot Encoding
 - ❑ Sparse way of representing
- ❑ Target Encoding
 - ❑ Replacing with the mean of the target
- ❑ Embeddings
 - ❑ Convert discrete values to a vector of continuous values

puppy	[0.9, 1.0, 0.0]
dog	[1.0, 0.2, 0.0]
kitten	[0.0, 1.0, 0.9]
cat	[0.0, 0.2, 1.0]

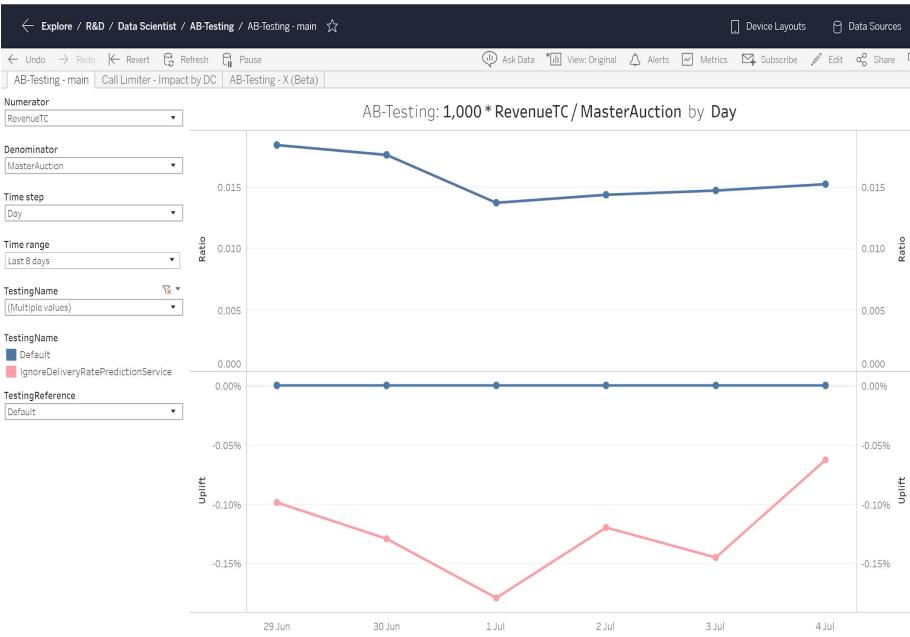
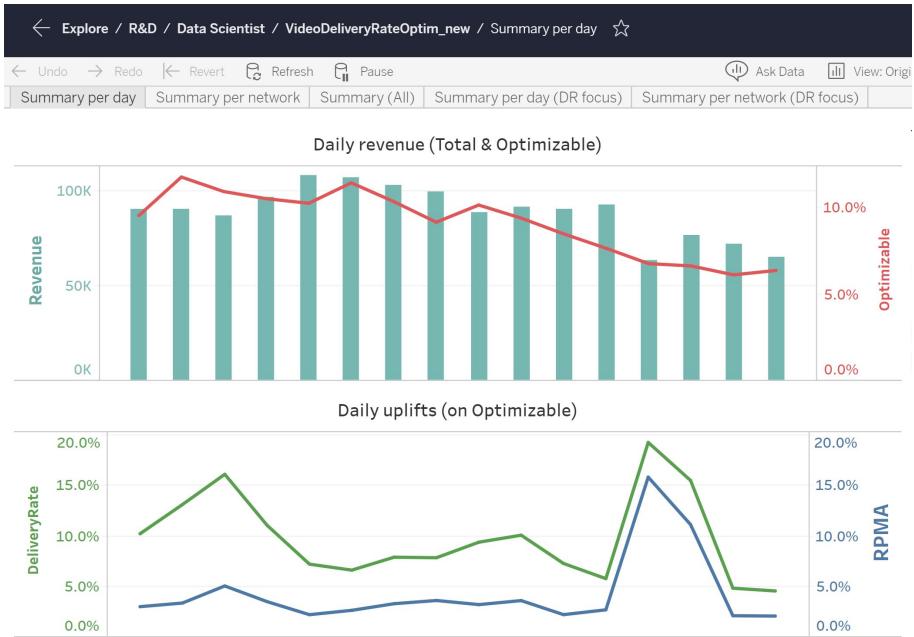




Video Optim - Next Steps

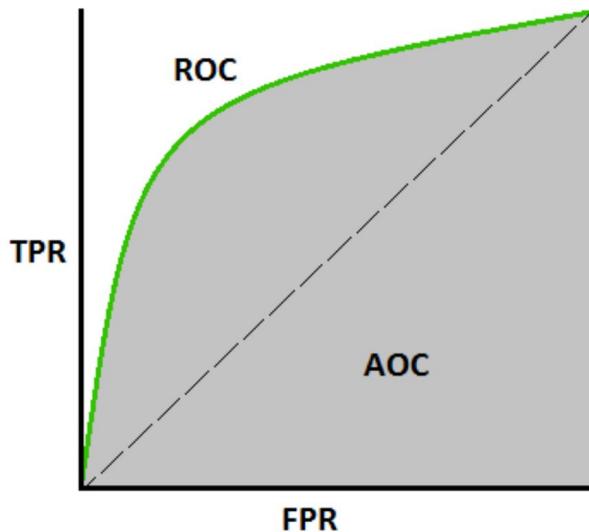
- Monitoring revenue uplift
 - Investigate results, find where DROS has big uplift / downlift
- Comparing model A and model B
- Anomaly detection workflow
 - Find the network responsible
 - Get the responsible dimensions (CreativeHashCode, Referrer)
- Exploration of VAST data + add important VAST dimensions to the model
- Change weighting strategy of the model
 - *Revenue and OptimizableRatio*
 - Focus where the impact is high
- Implementing continuous learning for deep learning model
- Industrialization of Deep Learning Model

Tableau Dashboards



Area Under the Curve

ROC (Receiver Operating Characteristics)



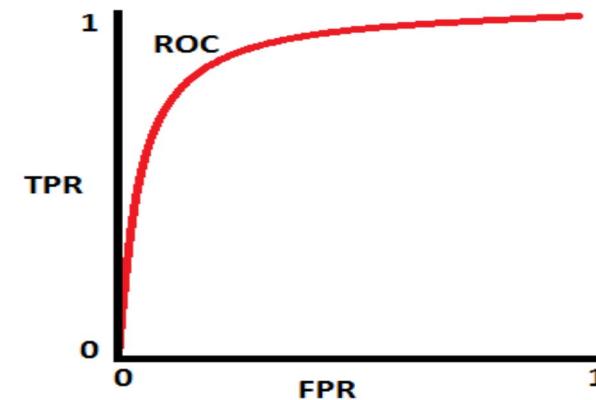
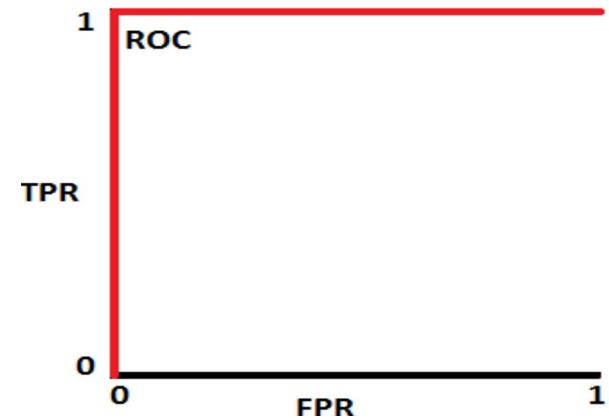
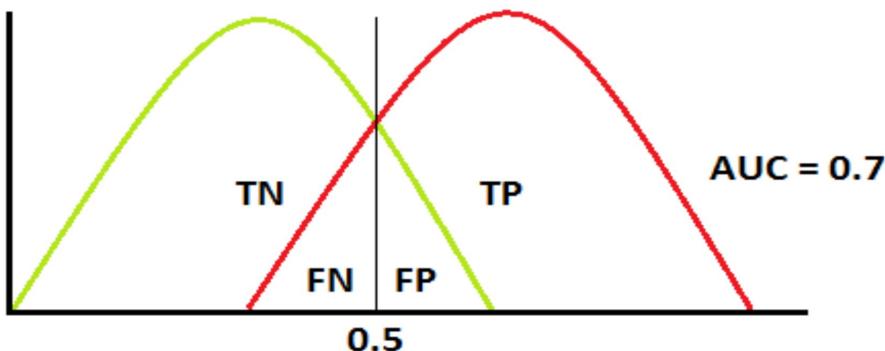
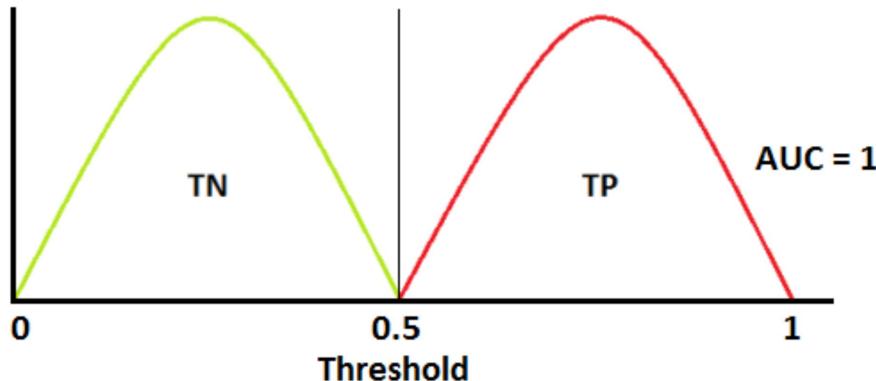
$$\text{TPR / Recall / Sensitivity} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{FPR} = 1 - \text{Specificity} = \frac{\text{FP}}{\text{TN} + \text{FP}}$$



ROC

curve of probability





XGBOOST LOSS FUNCTION

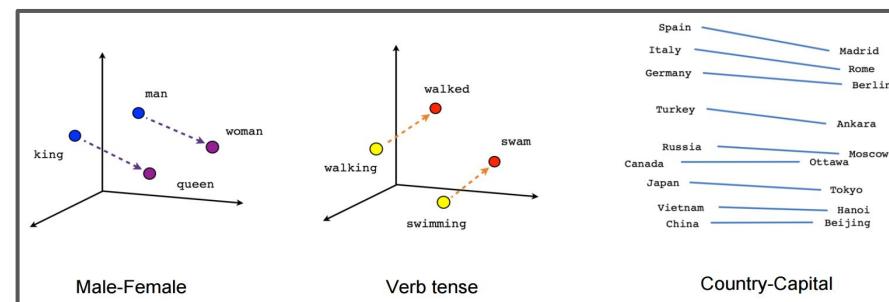
$$\left[\sum_{i=1}^n L(y_i, p_i) \right] + \frac{1}{2} \lambda O_{value}^2$$

PoC of a new model with Deep Learning

Handle Categorical Data with Deep Learning

- ❑ One Hot Encoding
 - ❑ Sparse way of representing
- ❑ Target Encoding
 - ❑ Replacing with the mean of the target
- ❑ Embeddings
 - ❑ Convert discrete values to a vector of continuous values

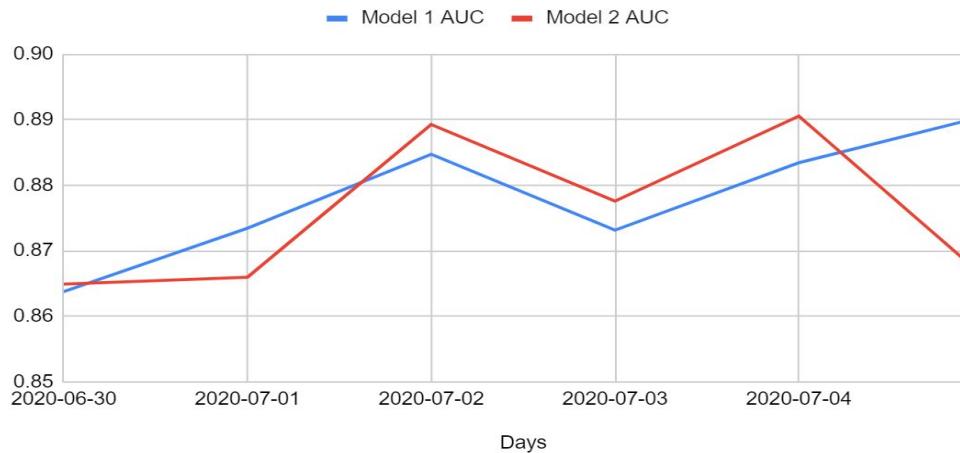
puppy	[0.9, 1.0, 0.0]
dog	[1.0, 0.2, 0.0]
kitten	[0.0, 1.0, 0.9]
cat	[0.0, 0.2, 1.0]



Model 1 Log Loss vs Model 2 Log Loss



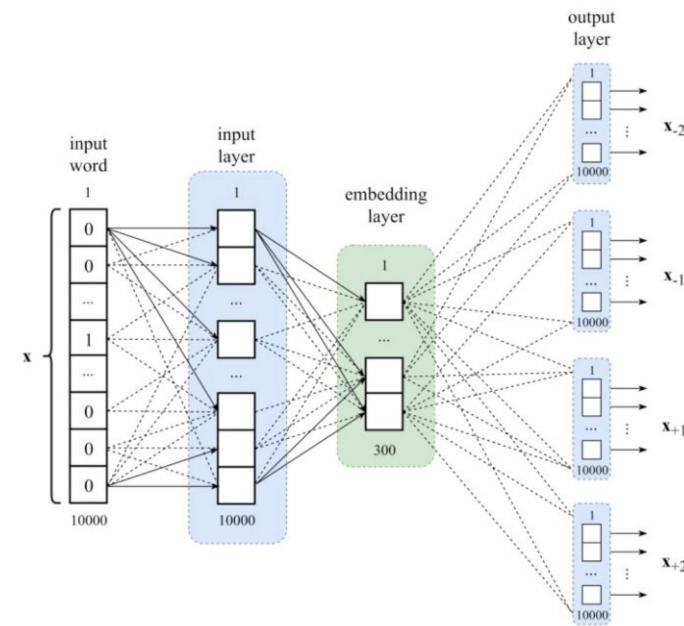
Model 1 AUC and vs Model 2 AUC



smart



Embedding Layer



smart

Moving all experiments to Google Cloud Platform

GCP

- ❑ Allocate resource as much as needed
- ❑ Flexible
- ❑ Easy to use



Google AI Platform

- ❑ Run ML code on top of TensorFlow, Scikit-learn and XGBoost
- ❑ Build ML pipelines



BigQuery

- ❑ Serverless, highly scalable, and cost-effective cloud data warehouse

smart

Switching Prediction Target from Success to Failure

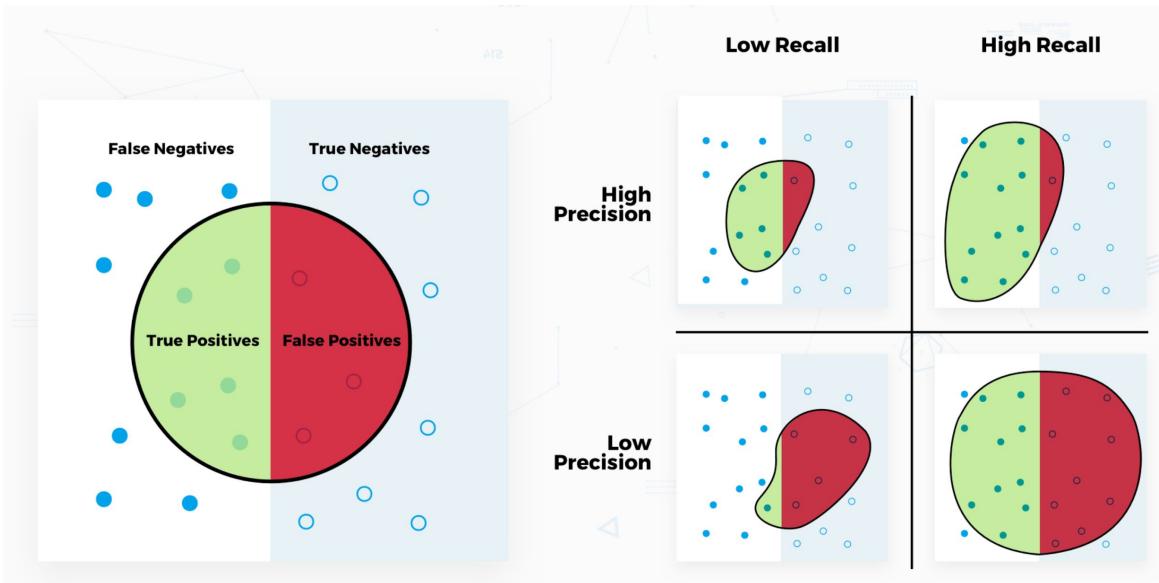
Before

Predicting Success



After

Predicting Failure



smart

ITX4

12 features;

IsHeaderBidding;HeaderBiddingTypeId;NetworkId;
SiteId;PlatformId;InsertionScriptId;ImpressionType;
Inventory;PartnerId;CreativeHashcode;
Referrer;InputBundleId

ITX5

19 features;

IsHeaderBidding;HeaderBiddingTypeId;NetworkId;
SiteId;PlatformId;InsertionScriptId;ImpressionType;
Inventory;PartnerId;CreativeHashcode;Referrer;
InputBundleId;InsertionId;FormatId;SmartBuyerId;
VideoPlacementType;PageId;OsParentId;
BrowserParentId

Perfumation feature importance

CreativeHashCode	0.083310
Pageld	0.041163
Referrer	0.025063
SmartBuyerId	0.023372
Inventory	0.019067
InsertionId	0.018772
PlatformId	0.018174
NetworkId	0.016077
BrowserParentId	0.011999
OsParentId	0.011998
PartnerId	0.011587
Siteld	0.007714
InputBundleId	0.005112
InsertionScriptId	0.002813
IsHeaderBidding	0.002100
VideoPlacementType	0.001938
FormatId	0.001512
HeaderBiddingTypeId	0.001163

Gain feature importance

Pageld	0.159561
CreativeHashCode	0.118734
Referrer	0.098011
Siteld	0.091464
VideoPlacementType	0.083967
NetworkId	0.063915
FormatId	0.055364
SmartBuyerId	0.048019
IsHeaderBidding	0.039537
InputBundleId	0.038564
PartnerId	0.037207
Inventory	0.030380
HeaderBiddingTypeId	0.027295
ImpressionType	0.027182
InsertionId	0.024226
InsertionScriptId	0.021150
PlatformId	0.012486
BrowserParentId	0.012018
OsParentId	0.010920

Target Encoding

is applying additive smoothing. The intuition behind this approach is smoothing the

computed mean by the global mean. For instance, in our case, the target is the delivery

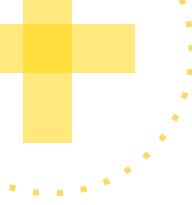
rate, and we are interested in encoding the networkId column. Assuming that there is a

new networkId that is contained by only three samples in the data set and it gives 0.9

of target mean for the new networkId whereas the target mean of all networkIds hover

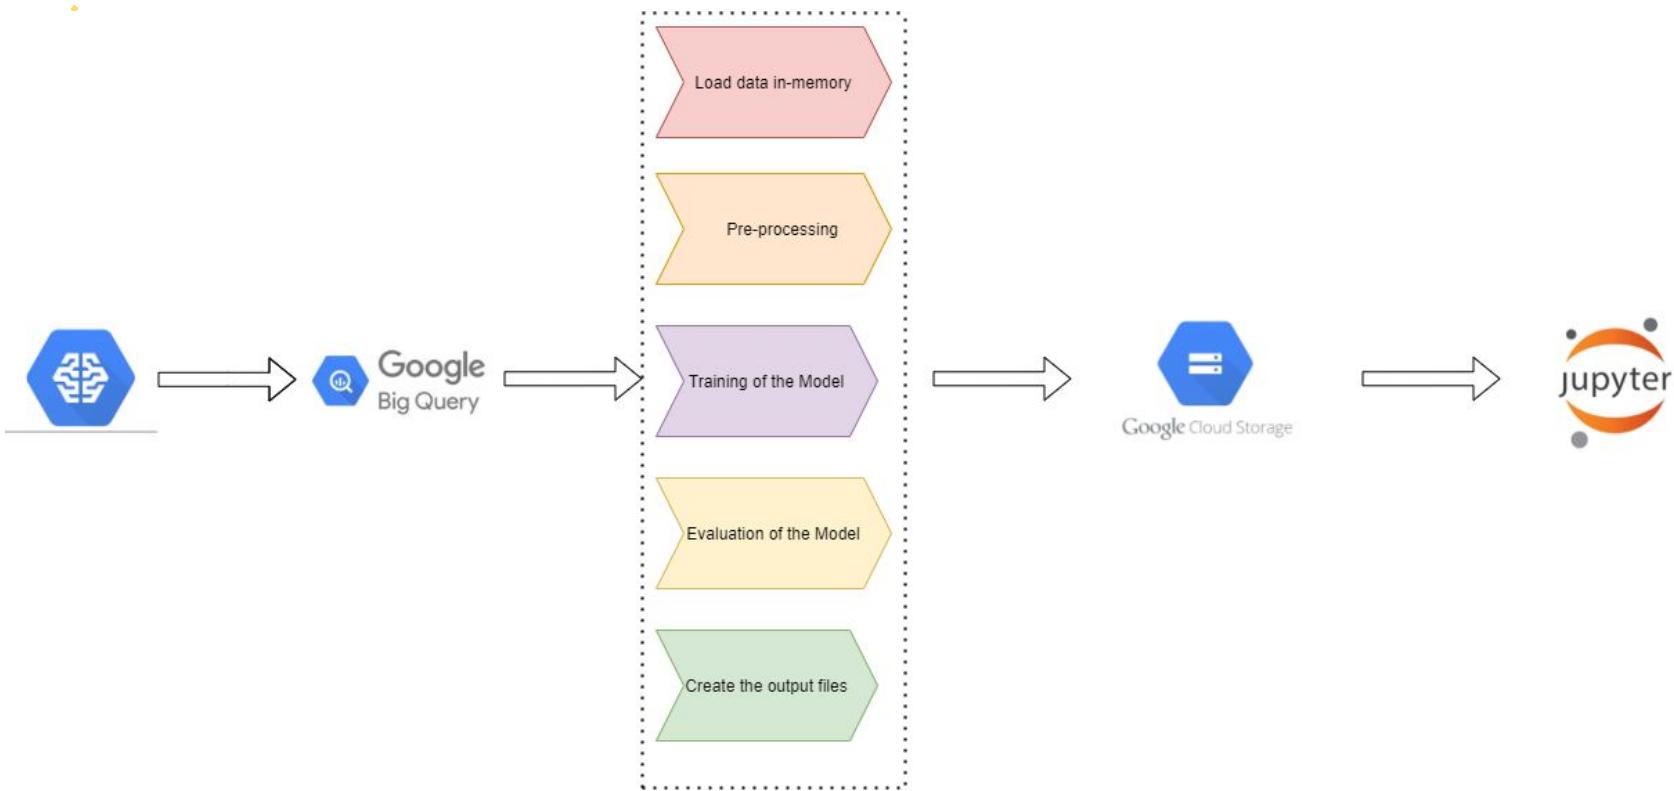
around 0.3. Therefore, it is better not to rely on only three values and to smooth the

average by including the average delivery rate over all the networkIds.

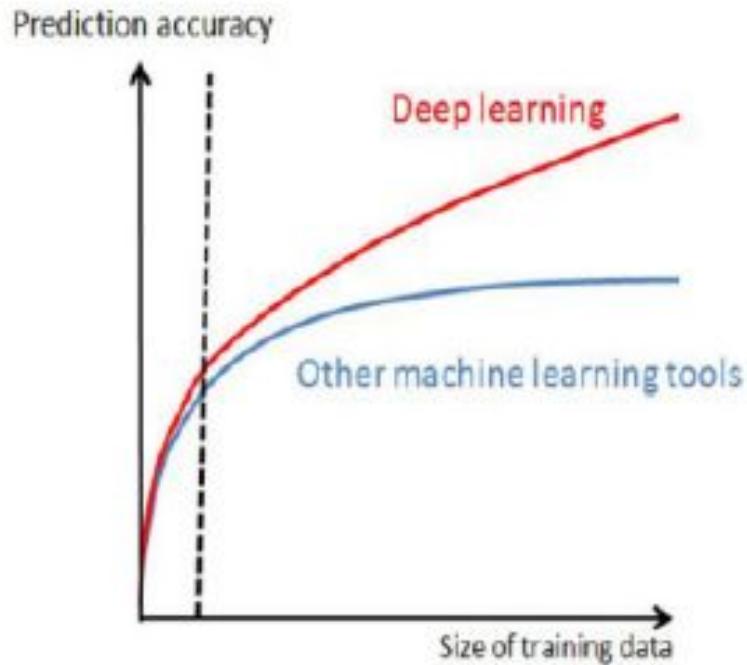


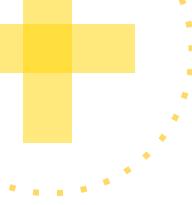
Embedding Layer

it limits the number of columns per category, and it is able to capture similarities between similar variables. It performs well in many benchmarks. However, it requires more tuning compared to the previous methods therefore it takes more time to implement, as explained in the next paragraph.



smart





XGBoost scalable and parallelizable

- ❖ Split Finding Algorithms(Basic Exact Greedy Algorithm)
- ❖ Approximate Algorithm:
 - XGBoost uses approximation on the split points
 - Quantile approach
 - Get the same accuracy as the exact greedy algorithm
 - Benefit of being distributable
 - ❖ Sparsity Aware Split Finding
 - ❖ Column Block for Parallel Learning
 - Block structure helps to optimize the computation complexity of split finding
 - Sort the data by column, in a compressed format
 - Different blocks can be distributed across machines
 - ❖ Cache-Aware Access(non-continuous)



TensorFlow Serving in Production

- ❖ Can serve multiple models, or multiple versions of the same model simultaneously
- ❖ Exposes both gRPC as well as HTTP/2 inference endpoints
- ❖ Allows deployment of new model versions without changing any client code
- ❖ Supports many *servables*: Tensorflow models, embeddings, vocabularies, feature transformations and even non-Tensorflow-based machine learning models